

Floating Point Arithmetics + Posits (Part 1)

by Uğur Çayoğlu

05. February 2020

STEINBUCH CENTRE FOR COMPUTING (SCC) &
INSTITUTE FOR METEOROLOGY AND CLIMATE RESEARCH (IMK-ASF)

Agenda




- **Floating Point Arithmetics (Part 1)**
- **Posits (Part 2)**

Take-home message of today!



$$0.1 + 0.2 \neq 0.3$$

- The correct answer is not of interest
 - **Accept life** and get over it!
 - “If life gives you lemons, ...”
 - Be consistent and **repeat mistakes!**
- Repeat mistakes **within one build**, across **multiple builds**, and across **multiple platforms**




Source_1

Structure and Format of Floating-Point Data



- Sign (1) + Exponent (8) + Mantissa (23) = IEEE Floating-Point Data (32)
 - Sign = Negative/Positive
 - Exponent = Power of two
 - Mantissa = Fraction ($1/2^{23} = 1/8,388,608 = \text{eps}$)
- $x = (-1)^s * 2^{(\text{exp}-127)} * (1+f)$




Source_2

Some Weird Properties of Floating-Point Data



- There are some special values (i.e. denormals):
 - +infinite, -infinite,
 - +zero, -zero,
 - NaN (there are actually 16,777,216 NaN values ($2 * \text{mantissa}$))



Source_2

Guarantees while using Floating-Point Arithmetics



These operations are determined and will produce **reproducible** results...

- Addition
- Subtraction
- Division
- Multiplication
- Square root computation

... if we use the same **rounding mode**, same **inputs**, same **global settings** and same destination **precision**

A Set of Problems Ordered by Your Likelihood of Experiencing it



1. Rounding modes
2. Composition
3. Per processor code
4. Precision modes
5. Denormals
6. Compiler differences / uninitialised data / transcendentals / square root estimates / conversions

Please check the trusted documentation of your programming language and compiler of choice for possibilities to prevent these problems

Please check your code before you blame floating point arithmetics...



- Check your **algorithms and data structures** before you blame floating point arithmetics about missing determinism
- Random number generator working properly?
- Flexible timing with thread scheduling?
- Simulation time variable?
- Code with undefined behaviour?



Source_3

Five Rounding modes



- Rounding mode for values in between two floating point values
- Runtime check necessary
- Thread bound
- C/C++: `#include <fenv.h>`

	11.5	12.5	- 11.5	- 12.5
TO_NEAREST_EVEN	12	12	-12	-12
TO_NEAREST_AWAY_ZERO	12	13	-12	-13
TOWARD_ZERO	11	12	-11	-12
PLUS_INFINITY	12	13	-11	-12
MINUS_INFINITY	11	12	-12	-13

Composition



- Optimization error
- Reordering of operations based on register/cache availability
- Throughput vs precision
- Use parenthesis for operation ordering
- C/C++: fp/fast vs fp/precise

$$a + b + c = ?$$

$$a + (b + c) \neq (a + b) + c$$

Precision Modes



- Different precisions for intermediate representations
- In most common architectures (i.e. Intel, AMD) this is 24-, 53-, 64-bit for representing mantissa
- Depends on the feature set one uses
 - C78: Floats are doubles where fraction part is filled with zeros
 - C++98: Only if one of the operands is a double
 - C99: “Evaluation type may be wider than semantic type”

... basically it is up to the compiler and the compilation flags used.

Per processor code



- Depending on the feature optimizations available the results might differ
- SSE = 32bit intermediate precision
- SSE2 = 64bit intermediate precision
- Special instructions
 - i.e. fmadd

$$\text{fmadd}(a, b, c) = a * b + c$$

$$a * b + c * d = ?$$

$$t = c * d$$

$$\text{fmadd}(a, b, t) \neq a * b + c * d$$

Denormals




- Special values from 0 to 1
- $x = (-1)^s * 2^{(\text{exp}-127)} * (1+f)$
 - Smallest possible value is for $f = 0$, $\text{exp} = 1$ (since $\text{exp}=0$ is +zero)
 - $2^{-126} = 1.175\text{E-}38$ (realmin)
 - $2^{-126} * 2^{-23} = 1.4013\text{e-}45$ (eps*realmin)
- All values between realmin and eps*realmin are called **denormals**
- Computations between these values are not hardware supported
- There is a flag to turn calculation of denormals off
 - All values below realmin will be set to 0
- Turn off if you want speed up, but results will change (> better don't)

Compiler differences / transcendentals / square root estimates / conversions



- Compile time/Runtime calculations (e.g. sin, tan, cos)
- Compiler: gcc, clang, visual c++, ...
- Square root estimates
 - There are square root estimates that are implemented by hardware manufacturers (rcpss, rcpps, rsqrtps, rsqrtss)
- Transcendentals like sin, cos, tan, pi, e can be hardware supported and might differ based on manufacturer
- Conversions e.g. printing commands might differ in output depending on operation order

That's all folks. Thank you...



<https://0.30000000000000004.com/>

Source_4

Resources



Main source for presentation: <http://randomascii.wordpress.com>

What Every Computer Scientist Should Know About Floating-Point Arithmetic

https://docs.oracle.com/cd/E19957-01/806-3568/ncg_goldberg.html

Float Converter <https://www.h-schmidt.net/FloatConverter/IEEE754.html>

Pics:

1. <https://i.imgur.com/Sgnwoln.jpg>
2. https://fabiansanglard.net/floating_point_visually_explained/floating_point_window_pi.svg
3. <https://static.inspiremore.com/wp-content/uploads/2017/01/24092529/Screen-Shot-2017-01-24-at-3.10.33-PM.png>
4. https://twitter.com/calvinn_hobbes/status/528720685771026432